

Lecture 2: Intro to pandas

TECH2: Introduction to Programming, Data, and Information Technology

Richard Foltyn

NHH Norwegian School of Economics

October 7, 2026

See GitHub repository for notebooks and data:

<https://github.com/richardfoltyn/TECH2-H26>

Contents

1	Introduction to pandas	1
1.1	Motivation	1
1.2	Creating pandas data structures	2
1.3	Importing data	4
1.4	Viewing data	6
1.5	Indexing	8
1.6	Working with time series data	17
1.7	List of functions used in this lecture	19

1 Introduction to pandas

1.1 Motivation

So far, we have encountered built-in Python containers (`list`, `tuple`, `dict`) and NumPy arrays as the only way to store data. However, while NumPy arrays are great for storing *homogeneous* data without any particular structure, they are somewhat limited when we want to use them for data analysis.

For example, we usually want to process data sets with

1. several variables;
2. multiple observations, which need not be identical across variables (some values may be missing);
3. non-homogeneous data types: for example, names need to be stored as strings, birthdays as dates, and income as a floating-point number.

While NumPy can in principle handle such situations, it puts all the burden on the user. Most users would prefer to not have to deal with such low-level details.

pandas was created to offer more versatile data structures that are straightforward to use for storing, manipulating, and analyzing heterogeneous data:

1. Data is organized in *variables* and *observations*, similar to econometrics programs such as Stata, or R using `data.frame` objects.
2. Each variable is permitted to have a *different* data type.

3. We can use *labels* to select observations instead of having to use a linear numerical index as with NumPy.

We could, for example, index a data set using National Insurance Numbers or time stamps for time series data.

4. pandas offers many convenient data aggregation and reduction routines that can be applied to subsets of data.

For example, we can easily group observations by city and compute average incomes.

5. pandas also offers many convenient data import / export functions that go beyond what's in NumPy.

Should we be using pandas at all times, then? No!

- For low-level tasks where performance is essential, use NumPy.
- For homogeneous data without any particular data structure, use NumPy.
- On the other hand, if data is heterogeneous, needs to be imported from an external data source and cleaned or transformed before performing computations, use pandas.

There are numerous tutorials on pandas on the internet. Useful additional material includes:

- The official [user guide](#).
- The official [pandas cheat sheet](#) which nicely illustrates the most frequently used operations.
- The official [API reference](#) with details on every pandas object and function.
- There are numerous tutorials (including videos) available on the internet. See [here](#) for a list.

1.2 Creating pandas data structures

pandas has two main data structures:

1. [Series](#) represents observations of a *single* variable.
2. [DataFrame](#) is a container for *several* variables. You can think of each individual column of a DataFrame as a Series, and each row represents one observation.

The easiest way to create a Series or DataFrame is to create them from pre-existing data.

To access pandas data structures and routines, we need to import them first. The near-universal convention is to make pandas available using the name `pd`:

```
import pandas as pd
```

Example: Create Series from 1-dimensional NumPy array

```
[1]: import numpy as np
import pandas as pd # universal convention: import using pd

# Create array of integers from 5 to 9
data = np.arange(5, 10)

# Create pandas Series from 1-dimensional NumPy array
pd.Series(data)
```

```
[1]: 0    5
      1    6
      2    7
      3    8
      4    9
      dtype: int64
```

Example: Create DataFrame from NumPy array

We can also create a DataFrame from a 2-dimensional NumPy array. To this end, we first create a 1-dimensional array of increasing integers from 0 to 14 and then use the `reshape()` method to reshape it to a 2-dimensional array.

```
[2]: # Create 5-by-3 matrix of data
data = np.arange(15).reshape((-1, 3))

# Define variable (or column) names
varnames = ['A', 'B', 'C']

# Create pandas DataFrame from matrix (2-dimensional NumPy array)
pd.DataFrame(data, columns=varnames)
```

```
[2]:   A  B  C
0  0  1  2
1  3  4  5
2  6  7  8
3  9 10 11
4 12 13 14
```

This code creates a DataFrame of three variables called A, B, and C with 5 observations each.

Example: Create DataFrame from dictionary

Alternatively, we can create a DataFrame from non-homogeneous data as follows:

```
[3]: # Names (strings)
names = ['Alice', 'Bob']

# Birth dates (datetime objects)
bdates = pd.to_datetime(['1985-01-01', '1997-05-12'])

# Incomes (floats)
incomes = np.array([60000.0, np.nan]) # code missing income as NaN

# create DataFrame from dictionary
pd.DataFrame({'Name': names, 'Birthdate': bdates, 'Income': incomes})
```

```
[3]:   Name  Birthdate  Income
0  Alice 1985-01-01  60000.0
1   Bob 1997-05-12     NaN
```

If data types differ across columns, as in the above example, it is often convenient to create the DataFrame by passing a dictionary as an argument. Each key represents a column name and each corresponding value contains the data for that variable.

Your turn

Create a pandas Series which contains the characters 'a', 'b', and 'c'.

1.3 Importing data

1.3.1 Digression: Loading data with NumPy & its limitations

Before turning to pandas, let's examine how one could import data using a NumPy function. This also highlights the limitations of this approach.

We often use files that store data as text files containing character-separated values (CSV) since virtually any application supports this data format. The most important NumPy functions to read text data are:

- `np.loadtxt()`: load data from a text file.
- `np.genfromtxt()`: load data from a text file and handle missing data.

There are a few other input/output functions in NumPy, for example to write arrays as raw binary data. We won't cover them here, but you can find them in the [official documentation](#).

Example: Load character-separated text data

Consider the following tabular data from [FRED](#) stored in the file [FRED_annual.csv](#) where the first two rows look as follows:

Year	GDP	CPI	UNRATE	FEDFUNDS	INFLATION
1954	2877.7	26.9	5.6	1.0	
1955	3083.0	26.8	4.4	1.8	-0.4

Note that the inflation column has a missing value for the year 1954.

These data are stored as character-separated values (CSV). To load this CSV file as a NumPy array, we use `loadtxt()`.

It is convenient to define the path to the data/ directory that can point either to the local directory or to the data/ directory on GitHub. We can then use this path to load the data file.

```
[4]: # Use files in the local data/ directory
DATA_PATH = '../..data'

# Alternatively, load data directly from GitHub
# DATA_PATH = 'https://raw.githubusercontent.com/richardfoltn/TECH2-H26/main/data'
```

```
[5]: import numpy as np

# Path to CSV file
file = f'{DATA_PATH}/FRED/FRED_annual.csv'

# load CSV, skip header row and first row with missing data
data = np.loadtxt(file, skiprows=2, delimiter=',')

data[:2] # Display first two rows
```

```
[5]: array([[ 1.9550e+03,  3.0830e+03,  2.6800e+01,  4.4000e+00,  1.8000e+00,
            -4.0000e-01],
            [ 1.9560e+03,  3.1488e+03,  2.7200e+01,  4.1000e+00,  2.7000e+00,
             1.5000e+00]])
```

The default settings work for many simple whitespace-delimited text files. However, we'll occasionally want to specify the following arguments to override the defaults:

- `delimiter`: Character used to separate individual fields (default: whitespace).
- `skiprows=n`: Skip the first `n` rows. For example, if the CSV file contains a header with variable names, `skiprows=1` needs to be specified as NumPy by default cannot process these names.

- **encoding:** Set the character encoding of the input data. This is usually not needed, but can be required to import data when the file's encoding differs from the default, for example when importing text encoded using Latin-1 instead of UTF-8.

While `loadtxt()` is simple to use, it quickly reaches its limits with more complex data sets. For example, when we try to load the FRED data set including the first data row, we get the following error:

```
[6]: # Attempt to load CSV
data = np.loadtxt(file, skiprows=1, delimiter=',')
```

```
ValueError: could not convert string '' to float64 at row 0, column 6.
```

This code fails because `loadtxt()` does not support files with missing values. One can use the more flexible function `np.genfromtxt()` which allows us to parse files with missing values:

```
[7]: # Load CSV file using genfromtxt() instead of loadtxt()
data = np.genfromtxt(file, skip_header=1, delimiter=',')

# Display first rows
data[:1]
```

```
[7]: array([[1.9540e+03, 2.8777e+03, 2.6900e+01, 5.6000e+00, 1.0000e+00,
            nan]])
```

However, it is usually not worthwhile to figure out how to load complex data with NumPy as this is much easier with pandas.

1.3.2 Loading data with pandas

pandas's input/output routines are more powerful than those implemented in NumPy:

- They support reading and writing numerous file formats.
- They support heterogeneous data without having to specify the data type in advance.
- They gracefully handle missing values.

For these reasons, it is often preferable to directly use pandas to process data instead of NumPy.

The most important functions are:

- `read_csv()`, `to_csv()`: Read or write CSV text files.
- `read_fwf()`: Read data with fixed field widths, i.e., text data that does not use delimiters to separate fields.
- `read_excel()`, `to_excel()`: Read or write Excel spreadsheets.
- `read_stata()`, `to_stata()`: Read or write Stata's `.dta` files.
- `read_json()`, `to_json()`: Read or write JSON files (lightweight machine-readable text files, often produced by LLMs)

For a complete list of I/O routines, see the [official documentation](#).

To illustrate, we repeat the above examples using pandas's `read_csv()`:

```
[8]: import pandas as pd

# Path to CSV file
file = f'{DATA_PATH}/FRED/FRED_annual.csv'

# Load CSV file using pandas's read_csv() function
```

```
df = pd.read_csv(file, sep=',')

# Display the first 2 rows of data
df.head(2)
```

```
[8]:
```

	Year	GDP	CPI	UNRATE	FEDFUNDS	INFLATION
0	1954	2877.7	26.9	5.6	1.0	NaN
1	1955	3083.0	26.8	4.4	1.8	-0.4

Your turn

Use the pandas functions listed above to import data from the following files located in the data folder:

1. `titanic.csv`
2. `FRED/FRED_annual.xlsx`

To load Excel files, you need to have the package `openpyxl` installed.

1.4 Viewing data

With large data sets, you hardly ever want to print the entire `DataFrame`. pandas by default limits the amount of data shown. You can use the `head()` and `tail()` methods to explicitly display a specific number of rows from the top or the end of a `DataFrame`.

To illustrate, we use a data set of passengers on board the Titanic's maiden voyage stored in `titanic.csv` which contains the following columns:

1. `PassengerId`
2. `Survived`: indicator whether the person survived
3. `Pclass`: accommodation class (first, second, third)
4. `Name`: Name of passenger (last name, first name)
5. `Sex`: male or female
6. `Age`
7. `Ticket`: Ticket number
8. `Fare`: Fare in pounds
9. `Cabin`: Deck + cabin number
10. `Embarked`: Port at which passenger embarked: C - Cherbourg, Q - Queenstown, S - Southampton

We can read in the data stored in the file `titanic.csv` like this:

```
[9]: import pandas as pd

# Path to CSV file
file = f'{DATA_PATH}/titanic.csv'

# Load sample data set of Titanic passengers. Individual fields are separated
# using a comma, which is the default.
df = pd.read_csv(file, sep=',')
```

We can now display the first and last three rows:

```
[10]: df.head(3) # show first three rows
```

```
[10]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	

		Name	Sex	Age	\
0		Braund, Mr. Owen Harris	male	22.0	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...		female	38.0	
2		Heikkinen, Miss Laina	female	26.0	

		Ticket	Fare	Cabin	Embarked
0	A/5	21171	7.2500	NaN	S
1	PC	17599	71.2833	C85	C
2	STON/O2.	3101282	7.9250	NaN	S

```
[11]: df.tail(3) # show last three rows
```

```
[11]:
```

	PassengerId	Survived	Pclass		Name	\
888	889	0	3	Johnston, Miss Catherine Helen	"Carrie"	
889	890	1	1	Behr, Mr. Karl Howell		
890	891	0	3	Dooley, Mr. Patrick		

	Sex	Age	Ticket	Fare	Cabin	Embarked
888	female	NaN	W./C. 6607	23.45	NaN	S
889	male	26.0	111369	30.00	C148	C
890	male	32.0	370376	7.75	NaN	Q

1.4.1 Setting pandas display options

It is possible to configure how much data pandas displays by default using `set_option()`. There is a large number of display options you can set; see [here](#) for the full list.

For our purposes, the following are the most important:

```
[12]: # Set the line width to 120 characters (default: 80)
pd.set_option('display.width', 120)
# Show up to 100 rows (default: 60)
pd.set_option('display.max_rows', 100)
```

1.4.2 Descriptive statistics

To quickly compute some descriptive statistics for the *numerical* variables in the DataFrame, we use `describe()`:

```
[13]: df.describe()
```

```
[13]:
```

	PassengerId	Survived	Pclass	Age	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	32.204208
std	257.353842	0.486592	0.836071	14.526497	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	512.329200

Note that this automatically ignores the columns Name, Sex, Ticket, Cabin, and Embarked as they contain strings, and computing the mean, standard deviation, etc. of a string variable does not make sense.

For categorical data, we can use `value_counts()` to tabulate the number of observations for each unique value of a variable. For example, the following code tabulates passengers by sex:

```
[14]: df['Sex'].value_counts()
```

```
[14]: Sex
male      577
female    314
Name: count, dtype: int64
```

Lastly, to see low-level information about the data type used in each column and the number of nonmissing observations, we call `info()`:

```
[15]: df.info(show_counts=True)
```

```
<class 'pandas.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    str
4   Sex          891 non-null    str
5   Age          714 non-null    float64
6   Ticket       891 non-null    str
7   Fare         891 non-null    float64
8   Cabin        204 non-null    str
9   Embarked     889 non-null    str
dtypes: float64(2), int64(3), str(5)
memory usage: 69.7 KB
```

Many pandas reductions, such as computing the mean, ignore missing values by default, whereas arithmetic generally propagates missing values. The `info()` output above reports the number of nonmissing observations in each column. For example, the Age column has several missing values, so its reported Non-Null count is lower than the total number of rows.

Starting with pandas 3.0, columns containing strings use the dedicated `str` data type by default, as shown above. These columns can contain strings and missing values, but cannot contain other types of values.

Your turn

Using the Titanic data set:

1. Tabulate the number of passengers by boarding port (variable Embarked).
2. Count the observations with a missing value for Embarked.

Hint: `value_counts()` excludes missing values by default. Use `dropna=False` to include them in the table, or use `isna().sum()` to count them separately.

1.5 Indexing

pandas supports two types of indexing:

1. Indexing by position. This is basically identical to the indexing of other Python and NumPy containers.
2. Indexing by label, i.e., by the values assigned to the row or column index. These labels need not be integers in increasing order, as is the case for NumPy. We will see how to assign labels below.

pandas indexing is performed either by using brackets [], or by using .loc[] for label indexing, or .iloc[] for positional indexing.

Indexing via [] can be somewhat confusing:

- specifying df['name'] returns the column name as a Series object.
- On the other hand, specifying a range such as df[5:10] returns the *rows* associated with the *positions* 5,...,9.

Example: Selecting columns

```
[16]: import pandas as pd

# Set option to limit the number of rows displayed
pd.set_option('display.max_rows', 7)

# Load sample data of Titanic passengers
df = pd.read_csv(f'{DATA_PATH}/titanic.csv')
df['Name'] # select a single column
```

```
[16]: 0          Braund, Mr. Owen Harris
1    Cumings, Mrs. John Bradley (Florence Briggs Th...
2          Heikkinen, Miss Laina
...
888    Johnston, Miss Catherine Helen "Carrie"
889          Behr, Mr. Karl Howell
890    Dooley, Mr. Patrick
Name: Name, Length: 891, dtype: str
```

```
[17]: df[['Name', 'Sex']] # select multiple columns using a list
```

```
[17]:
```

	Name	Sex
0	Braund, Mr. Owen Harris	male
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female
2	Heikkinen, Miss Laina	female
..	...	...
888	Johnston, Miss Catherine Helen "Carrie"	female
889	Behr, Mr. Karl Howell	male
890	Dooley, Mr. Patrick	male

[891 rows x 2 columns]

To select multiple columns, pass a list of column names as shown above. Other suitable list-like objects, such as NumPy arrays, also work.

Important: A tuple is treated as a single column label, so it does not select multiple columns in this example.

Example: Selecting rows by position

To return the rows at positions 1, 2, and 3 we use

```
[18]: df[1:4]
```

```
[18]:
```

PassengerId	Survived	Pclass
Name	Sex	Age
1	2	1
Th...	female	38.0
2	3	1
Laina	female	26.0
3	4	1
Peel)	female	35.0

	Fare	Cabin	Embarked
1	71.2833	C85	C
2	7.9250	NaN	S
3	53.1000	C123	S

Positional indexing in pandas follows the Python convention: positions are 0-based, and the endpoint of a positional slice is excluded. Index labels need not be integers, and label slices include both endpoints.

1.5.1 Creating and manipulating indices

pandas uses *labels* to index and align data. These can be integer values starting at 0 with increments of 1 for each additional element, which is the default, but they need not be. The three main methods to create/manipulate indices are:

1. Create a new Series or DataFrame object with a custom index using the `index` argument.
2. `set_index(keys=['column1', ...])` uses the values of `column1` and optionally additional columns as indices, discarding the current index.
3. `reset_index()` resets the index to its default value, a sequence of increasing integers starting at 0.

Creating custom indices

First, consider the following code which creates a Series with three elements `[10, 20, 30]` using the default index `[0, 1, 2]`:

```
[19]: import pandas as pd

# Create Series with default integer index
pd.Series([10, 20, 30])
```

```
[19]: 0    10
      1    20
      2    30
      dtype: int64
```

We can use the `index` argument to specify a custom index, for example one containing the lowercase characters `a`, `b`, and `c` as follows:

```
[20]: # Create Series with custom index [a, b, c]
pd.Series([10, 20, 30], index=['a', 'b', 'c'])
```

```
[20]: a    10
      b    20
      c    30
      dtype: int64
```

Manipulating indices

To modify the index of an *existing* DataFrame, we use the methods `set_index()` and `reset_index()`. By default, these return a new object and leave the original DataFrame unchanged. To modify the existing object directly, pass `inplace=True`. Alternatively, assign the returned object back to the original variable. A Series supports `reset_index()`, but does not have a `set_index()` method.

For example, we can replace the row index and use the Roman lowercase characters `a`, `b`, `c`, ... as labels instead of integers:

```
[21]: # Create DataFrame with 2 columns
df = pd.DataFrame({'A': [10, 20, 30], 'B': ['a', 'b', 'c']})

df
```

```
[21]:      A  B
0   10  a
1   20  b
2   30  c
```

Since we did not specify any index, the default index [0, 1, ...] is used. We can use `set_index()` to set the index to the values from a column, for example column B:

```
[22]: # Use column 'B' as index, store result in new DataFrame
df2 = df.set_index('B')

# Display updated DataFrame
df2
```

```
[22]:      A
B
a   10
b   20
c   30
```

In pandas 3.0, objects returned by indexing or methods such as `set_index()` behave as independent copies when modified. This is called *Copy-on-Write*: pandas may initially share the underlying data and copy it only when needed to prevent changes to one object from affecting another.

In the previous example, only `df2` uses column B as the index, whereas the original `df` remains unchanged:

```
[23]: df
```

```
[23]:      A  B
0   10  a
1   20  b
2   30  c
```

We can use the `inplace=True` argument to `set_index()` to update the index of the existing DataFrame:

```
[24]: # Set index in-place, i.e., df is modified
df.set_index('B', inplace=True)

df
```

```
[24]:      A
B
a   10
b   20
c   30
```

`set_index()` and `reset_index()` return `None` when called with `inplace=True`, so assigning their return value to a variable would replace that variable's value with `None`.

We can now use these new labels to select records in the DataFrame:

```
[25]: # print first 2 rows using labels
df['a':'b'] # This is the same as df[:2]
```

```
[25]:      A
B
```

```
a 10
b 20
```

Note that when specifying a range in terms of labels, the last element *is* included! Hence the row with index b in the above example is shown.

We can reset the index to its default integer values using the `reset_index()` method:

```
[26]: # Reset index labels to default value (integers 0, 1, 2, ...) and print
      # first three rows
      df.reset_index(drop=True).head(3)
```

```
[26]:      A
0    10
1    20
2    30
```

The `drop=True` argument tells pandas to throw away the old index values instead of storing them as a column of the resulting DataFrame.

Your turn

Read in the following data files from the `data/FRED` folder and manipulate the DataFrame index:

1. Read in the file `FRED_annual.csv` and set the column `Year` as the index.
2. Read in the file `FRED_monthly.csv` and set the columns `Year` and `Month` as the index.

Perform the tasks using `inplace=False` and `inplace=True`. What's the difference? Restore the original (default) index after you are done.

1.5.2 Selecting elements

To more clearly distinguish between selection by label and by position, pandas provides the `.loc[]` and `.iloc[]` methods of indexing. To make your intention obvious, you should therefore adhere to the following rules:

1. Use `df['name']` to select *columns*.
2. Use `.loc[]` to select by label.
3. Use `.iloc[]` to select by position.

Selection by label

To illustrate, using `.loc[]` unambiguously indexes by label. First we create a demo data set with 3 columns and 5 rows:

```
[27]: # Create demo data with 3 columns and 5 rows

      # Column labels
      columns = ['X', 'Y', 'Z']
      # Row labels
      rows = ['a', 'b', 'c', 'd', 'e']

      values = np.arange(len(rows))

      # Create data dictionary using a list comprehension
      data = {col: [f'{col}{val}' for val in values] for col in columns}

      # Create DataFrame from dictionary
```

```
df = pd.DataFrame(data, index=rows)

# Display DataFrame
df
```

```
[27]:
```

	X	Y	Z
a	X0	Y0	Z0
b	X1	Y1	Z1
c	X2	Y2	Z2
d	X3	Y3	Z3
e	X4	Y4	Z4

We now use `.loc[]` to select rows and columns by label. For example, to select the single row corresponding to the label 'b' we specify:

```
[28]: df.loc['b']
```

```
[28]:
```

	X	Y	Z
b	X1	Y1	Z1

Name: b, dtype: str

To select a specific element by row and column labels, we proceed as follows:

```
[29]: df.loc['b', 'X']
```

```
[29]: 'X1'
```

If you want to select multiple rows or columns, you can also pass lists as arguments to `.loc[]`:

```
[30]: # Select rows 'b' and 'd', and columns 'X' and 'Z'
df.loc[['b', 'd'], ['X', 'Z']]
```

```
[30]:
```

	X	Z
b	X1	Z1
d	X3	Z3

Finally, it is possible to perform slicing on rows or columns using the `:` notation which you might have encountered when slicing lists, tuples, or NumPy arrays:

```
[31]: # Select rows 'b' to 'e'
df.loc['b':'e']
```

```
[31]:
```

	X	Y	Z
b	X1	Y1	Z1
c	X2	Y2	Z2
d	X3	Y3	Z3
e	X4	Y4	Z4

With `.loc[]` we can even perform slicing on column names, which is not possible with the simpler `df[]` syntax:

```
[32]: # Select rows 'b' to 'd', and columns 'X' to 'Y'
df.loc['b':'d', 'X':'Y']
```

```
[32]:
```

	X	Y
b	X1	Y1
c	X2	Y2
d	X3	Y3

This includes all the columns between X and Y, where the latter is included since we are slicing by label.

Trying to use integer labels raises an error here because the index contains string labels (a, b, c, ...) rather than integers (0, 1, 2, ...).

```
[33]: df.loc[0:4]
```

```
TypeError: cannot do slice indexing on Index with these indexers [0] of type int
```

However, we can reset the index to its default value. Then the index labels are integers and coincide with their position, so that `.loc[]` works:

```
[34]: df = df.reset_index(drop=True) # reset index labels to integers,
      # drop original index
      df.loc[0:4]
```

```
[34]:
```

	X	Y	Z
0	X0	Y0	Z0
1	X1	Y1	Z1
2	X2	Y2	Z2
3	X3	Y3	Z3
4	X4	Y4	Z4

Again, the end point with label 4 is included because we are selecting by label.

Indexing via `.loc[]` supports a few more types of arguments; see the [official documentation](#) for details.

Selection by position

Conversely, if we want to select items exclusively by their position and ignore their labels, we use `.iloc[]`:

```
[35]: df.iloc[0:4, 0:2] # select first 4 rows, first 2 columns
```

```
[35]:
```

	X	Y
0	X0	Y0
1	X1	Y1
2	X2	Y2
3	X3	Y3

Again, `.iloc[]` supports a multitude of other arguments; see the [official documentation](#) for details.

Your turn

Using the DataFrame `df` we created above:

1. Compare `df.loc[0:3]` and `df.iloc[0:3]`. How many rows does each select? Predict the results before running the code.
2. Compare `df.loc[-1]` and `df.iloc[-1]`. What happens in each case, and why?

Boolean indexing

Similar to NumPy, pandas allows us to select a subset of rows in a Series or DataFrame if they satisfy some condition. The simplest use case is to create a column of Boolean values (True or False) as the result of a logical operation:

```
[36]: import pandas as pd

      # Read in Titanic passenger data
      df = pd.read_csv(f'{DATA_PATH}/titanic.csv')
```

```
# Check whether passenger embarked in Southampton
df['Embarked'] == 'S'
```

```
[36]: 0      True
      1     False
      2      True
      ...
      888    True
      889   False
      890   False
      Name: Embarked, Length: 891, dtype: bool
```

Such Boolean arrays can be used to select a subset of entries:

```
[37]: df.loc[df['Embarked'] == 'S', 'Name':'Age']
```

```
[37]:
```

	Name	Sex	Age
0	Braund, Mr. Owen Harris	male	22.0
2	Heikkinen, Miss Laina	female	26.0
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0
..	...	...	...
886	Montvila, Rev. Juozas	male	27.0
887	Graham, Miss Margaret Edith	female	19.0
888	Johnston, Miss Catherine Helen "Carrie"	female	NaN

[644 rows x 3 columns]

Boolean indexing also works directly with [] without having to specify .loc[], but a single df[] operation cannot select both rows and columns:

```
[38]: df[df['Embarked'] == 'S']
```

```
[38]:
```

	PassengerId	Survived	Pclass	Name
0	1	0	3	Braund, Mr. Owen Harris
male	22.0	A/5	21171	
2	3	1	3	Heikkinen, Miss Laina
female	26.0	STON/O2.	3101282	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)
female	35.0		113803	
..	...	...	...	...
...	...	...	...	...
886	887	0	2	Montvila, Rev. Juozas
male	27.0		211536	
887	888	1	1	Graham, Miss Margaret Edith
female	19.0		112053	
888	889	0	3	Johnston, Miss Catherine Helen "Carrie"
female	NaN	W./C.	6607	

	Fare	Cabin	Embarked
0	7.250	NaN	S
2	7.925	NaN	S
3	53.100	C123	S
..	...	...	...
886	13.000	NaN	S
887	30.000	B42	S
888	23.450	NaN	S

[644 rows x 10 columns]

Multiple conditions can be combined using the & (logical and) or | (logical or) operators:

```
[39]: # Select men who embarked in Southampton
df.loc[(df['Embarked'] == 'S') & (df['Sex'] == 'male'), ['Name', 'Embarked', 'Sex']]
```

```
[39]:
```

	Name	Embarked	Sex
0	Braund, Mr. Owen Harris	S	male
4	Allen, Mr. William Henry	S	male
6	McCarthy, Mr. Timothy J	S	male
..	...	...	...
883	Banfield, Mr. Frederick James	S	male
884	Sutehall, Mr. Henry Jr	S	male
886	Montvila, Rev. Juozas	S	male

[441 rows x 3 columns]

If we want to include rows where an observation takes on one of multiple values, the `isin()` method can be used:

```
[40]: # Select passengers who embarked in Southampton or Queenstown
df.loc[df['Embarked'].isin(('S', 'Q')), ['Name', 'Embarked']]
```

```
[40]:
```

	Name	Embarked
0	Braund, Mr. Owen Harris	S
2	Heikkinen, Miss Laina	S
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	S
..	...	...
887	Graham, Miss Margaret Edith	S
888	Johnston, Miss Catherine Helen "Carrie"	S
890	Dooley, Mr. Patrick	Q

[721 rows x 2 columns]

Finally, DataFrame implements a `query()` method which allows us to combine multiple conditions in a single string in an intuitive fashion. Column names can be used directly within this string to put restrictions on their values.

```
[41]: # Select passengers who embarked in Southampton and were above age 70
df.query('Embarked == "S" & Age > 70')
```

```
[41]:
```

	PassengerId	Survived	Pclass	Name	Sex
Age	Ticket	Fare	Cabin	Embarked	
63.0	631	1	1	Barkworth, Mr. Algernon Henry Wilson	male
80.0	27042	30.000	A23	S	
85.1	852	0	3	Svensson, Mr. Johan	male
74.0	347060	7.775	NaN	S	

Your turn

Load the Titanic passenger data set `data/titanic.csv` and select the following subsets of data:

1. Select the 10th to 20th (inclusive) rows of the DataFrame.
2. Using `query()`, select the sub-sample of female passengers aged 30 to 40. Display only the columns Name, Age, and Sex (in that order).
3. Repeat the last exercise without using `query()`.
4. Select all men who embarked in Queenstown or Cherbourg.

1.6 Working with time series data

In economics and finance, we frequently work with time series data, i.e., observations that are associated with a particular point in time (time stamp) or a time period. pandas offers comprehensive support for such data, in particular if the time stamp or time period is used as the index of a Series or DataFrame. This section presents a few of the most important concepts. See the official [documentation](#) for a comprehensive guide.

To illustrate, let's construct a set of daily data for the first three months of 2026, i.e., the period 2026-01-01 to 2026-03-31 using the `date_range()` function (we use the date format YYYY-MM-DD in this section, but pandas also supports other date formats).

```
[42]: import numpy as np
import pandas as pd

# Create sequence of dates from 2026-01-01 to 2026-03-31
# at daily frequency
index = pd.date_range(start='2026-01-01', end='2026-03-31', freq='D')

# Use date range as index for Series with some artificial data
data = pd.Series(np.arange(len(index)), index=index)

# Print first 5 observations
data.head(5)
```

```
[42]: 2026-01-01    0
      2026-01-02    1
      2026-01-03    2
      2026-01-04    3
      2026-01-05    4
      Freq: D, dtype: int64
```

1.6.1 Indexing with date/time indices

pandas implements several convenient ways to select observations associated with a particular date or a set of dates. For example, if we want to select one specific date, we can pass it as a string to `.loc[]`:

```
[43]: # Select single observation by date
data.loc['2026-01-01']
```

```
[43]: np.int64(0)
```

It is also possible to select a time period by passing a start and end point (where the end point is included, as usual with label-based indexing in pandas):

```
[44]: # Select first 5 days
data.loc['2026-01-01':'2026-01-05']
```

```
[44]: 2026-01-01    0
      2026-01-02    1
      2026-01-03    2
      2026-01-04    3
      2026-01-05    4
      Freq: D, dtype: int64
```

A particularly useful way to index time periods is to pass a partial index. For example, if we want to select all observations from January 2026, we could use the range `'2026-01-01': '2026-01-31'`, but it is much easier to specify the partial index `'2026-01'` instead which includes all observations from January.

```
[45]: # Select all observations from January 2026
data.loc['2026-01']
```

```
[45]: 2026-01-01    0
      2026-01-02    1
      2026-01-03    2
      ..
      2026-01-29   28
      2026-01-30   29
      2026-01-31   30
      Freq: D, Length: 31, dtype: int64
```

1.6.2 Lags, differences, and other useful transformations

When working with time series data, we often need to create lags or leads of a variable (e.g., if we want to include lagged values in a regression model). In pandas, this is done using `shift()` which, without a `freq` argument, shifts values by the desired number of rows (default: 1) while keeping the index unchanged. For example, invoking `shift(1)` creates lagged observations of a Series or each column in a DataFrame.

To shift date labels while preserving the values, specify a `freq` argument, such as `shift(1, freq='D')`.

```
[46]: # Lag observations by 1 period
data.shift(1).head(5)
```

```
[46]: 2026-01-01    NaN
      2026-01-02    0.0
      2026-01-03    1.0
      2026-01-04    2.0
      2026-01-05    3.0
      Freq: D, dtype: float64
```

We can use the `diff()` method to compute differences over a given number of periods:

```
[47]: # Compute difference between consecutive observations
data.diff(1).head(5)
```

```
[47]: 2026-01-01    NaN
      2026-01-02    1.0
      2026-01-03    1.0
      2026-01-04    1.0
      2026-01-05    1.0
      Freq: D, dtype: float64
```

Note that `diff()` is identical to manually computing the difference with the lagged value like this:

```
data - data.shift()
```

Additionally, we can use `pct_change()` which computes the fractional change (the relative difference) over a given number of periods (default: 1). Multiply the result by 100 to express the change as a percentage.

```
[48]: # Compute fractional change vs. previous period
data.pct_change().head(5)
```

```
[48]: 2026-01-01    NaN
      2026-01-02    inf
      2026-01-03    1.000000
      2026-01-04    0.500000
      2026-01-05    0.333333
      Freq: D, dtype: float64
```

This method is a shortcut for computing the fractional change as follows:

```
(data - data.shift()) / data.shift()
```

The first result is missing because there is no previous observation. The second is `inf` because the previous value is zero.

Your turn

Use the data from the `data/FRED` folder to perform the following tasks:

1. Read in the file `FRED_annual.csv` and set the column `Year` as the index.
2. Compute annual inflation as the percentage change of the consumer price index (column `CPI`). Use `pct_change()` and multiply by 100 to express inflation as a percentage.

1.7 List of functions used in this lecture

The following list contains the central functions covered in this lecture. Each function name is linked to the official API documentation which you can consult for more details regarding function arguments and return values. The [\[go to section\]](#) links to the relevant section in this notebook.

1.7.1 Creating pandas containers from existing data

- `pd.Series()` — create a Series containing observations of one variable [\[go to section\]](#)
- `pd.DataFrame()` — create a DataFrame containing observations of multiple variables [\[go to section\]](#)

1.7.2 Loading existing data

- `pd.read_csv()` — read CSV text file [\[go to section\]](#)
- `pd.read_excel()` — read Excel file [\[go to section\]](#)

1.7.3 Inspecting data

- `info()` — print information about observation count, columns, and data types [\[go to section\]](#)
- `head()` — print first few rows [\[go to section\]](#)
- `tail()` — print last few rows [\[go to section\]](#)
- `describe()` — print summary statistics for *numerical* data [\[go to section\]](#)
- `value_counts()` — tabulate observation counts for categorical data [\[go to section\]](#)

1.7.4 Manipulating the index

- `pd.Series(..., index=...)` — create Series with a custom index [\[go to section\]](#)
- `pd.DataFrame(..., index=...)` — create DataFrame with a custom index [\[go to section\]](#)
- `set_index()` — set existing columns as the index [\[go to section\]](#)
- `reset_index()` — reset index to default value [\[go to section\]](#)

1.7.5 Selecting observations

- `loc[]` — select elements by *label* (values of the index) [\[go to section\]](#)
- `iloc[]` — select elements by *position* [\[go to section\]](#)
- `isin()` — return whether elements are contained in a set of values [\[go to section\]](#)
- `query()` — select subset of observations according to some criteria [\[go to section\]](#)

1.7.6 Time series

- `date_range()` — create sequence of dates [\[go to section\]](#)
- `shift()` — shift values by the desired number of rows; shift date labels if `freq` is specified [\[go to section\]](#)
- `diff()` — calculate difference in levels vs. a previous/next element [\[go to section\]](#)
- `pct_change()` — calculate *relative* difference vs. a previous/next element [\[go to section\]](#)