

Workshop 4: Grouping and aggregation

TECH2: Introduction to Programming, Data, and Information Technology

Richard Foltyn
NHH Norwegian School of Economics

October 23, 2026

See GitHub repository for notebooks and data:
<https://github.com/richardfoltyn/TECH2-H26>

Contents

1	Exercise 1: From municipalities to counties	1
2	Exercise 2: House price levels and dispersion	2
3	Exercise 3: Determinants of house prices	2
4	Exercise 4: Inflation and unemployment in the US	3

Exercise 1: From municipalities to counties

Each Norwegian municipality (*kommune*) belongs to a county (*fylke*). In this exercise, you will use municipality-level data from Statistics Norway (SSB) to compare counties by their number of municipalities, population, and population density.

1. Load `municipalities.csv` from the `data/ssb/` folder. Keep only the columns `municipality`, `county`, `population`, and `land_area`, and display the first five rows.

Each row represents one municipality. The column `population` records its number of residents, and `land_area` records its land area in km^2 .

2. Group the data by county and count the municipalities in each county. Display the counts in descending order. Which county contains the most municipalities?

Hint: Applying `size()` to a grouped DataFrame counts the rows in each group.

3. Create a DataFrame named `df_counties` containing the total population and total land area of each county. Report the three counties with the largest populations.

Check that summing the county populations gives the same national total as summing the municipality populations.

4. Add a column named `density` to `df_counties`, containing each county's population density in residents per km^2 of land:

$$\text{Population density} = \frac{\text{Total population}}{\text{Total land area}}$$

Report the three counties with the highest population densities. Does the county with the largest population also have the highest population density?

Exercise 2: House price levels and dispersion

For this exercise, we're using data on around 3,000 observations of house prices and house characteristics from Ames, a small city in Iowa.

1. Load the Ames housing data set from `ames_houses.csv` located in the `data/` folder.

To familiarize yourself with the data, report the columns present in the DataFrame and then restrict the data to the columns `SalePrice` (price in thousands of US dollars) and `Neighborhood` (abbreviated name of the neighborhood).

Count the number of observations in each neighborhood. Drop all observations belonging to neighborhoods with fewer than 20 observations. Use this restricted data set for the remainder of the exercise.

Check that there are no observations with missing values in the final data set.

2. Compute the average house price (column `SalePrice`) by neighborhood (column `Neighborhood`). List the three neighborhoods with the highest average prices, for example by using `sort_values()`.

Create a bar chart showing the average sale price for all remaining neighborhoods in descending order.

Hint: You can create the bar chart by either using pandas's `DataFrame.plot.bar()` or Matplotlib's `bar()`.

3. You want to quantify the price dispersion in each neighborhood. To this end, compute the standard deviation of the sale price by neighborhood using `std()`. Which are the three neighborhoods with the most dispersed prices?
4. Create a scatter plot with the average house price on the x -axis and the standard deviation on the y -axis. Is there a relationship between the average house prices and their dispersion within a neighborhood?

Hint: You can create the scatter plot by either using pandas's `DataFrame.plot.scatter()` or Matplotlib's `scatter()`.

5. **[Advanced]** To compare neighborhoods with different average prices, express the standard deviation as a percentage of the average price:

$$\text{Relative price spread (\%)} = 100 \times \frac{\text{standard deviation}}{\text{average price}}$$

This is the **coefficient of variation** expressed as a percentage. For example, a standard deviation of 20 thousand USD and an average price of 100 thousand USD give a relative spread of 20%.

Calculate the relative price spread for each neighborhood. Create a scatter plot with the average house price on the x -axis and the relative price spread on the y -axis. Do neighborhoods with higher average prices also have more dispersed prices relative to their average? Compare the pattern with Part 4.

Exercise 3: Determinants of house prices

For this exercise, we're using data on around 3,000 observations of house prices and house characteristics from Ames, a small city in Iowa, to understand how house prices vary with selected house characteristics.

1. Load the Ames housing data set from `ames_houses.csv` located in the `data/` folder.
Keep only the columns `SalePrice` (price in thousands of US dollars), `LotArea` (lot area in m²), `YearBuilt`, and `Bedrooms` (number of bedrooms).
Restrict your data set to houses with one or more bedrooms and a lot area of at least 100m². Report the sample size after applying these restrictions.
2. Compute the average year in which a house was built (using the column `YearBuilt`).
Classify houses as “*newer*” if they were built after the sample’s average construction year, and “*older*” otherwise. Store this classification in a new column `Newer`, with 1 for newer houses and 0 for older houses.
What is the average year of construction within these two categories?
3. Create a new column `Rooms` which categorizes the number of `Bedrooms` into three groups: 1–2, 3, and 4 or more. Use the values 2, 3, and 4 to represent these groups. You can create these categories using boolean indexing, `np.where()`, pandas’s `where()`, `map()`, `clip()`, or some other way.
4. Compute the mean sale price and number of observations within each group formed by `Newer` and `Rooms` (for a total of 6 different categories).
By how much does the average price of newer houses with 4+ bedrooms exceed that of older houses with 1–2 bedrooms?
5. Create a figure with two subplots arranged in two columns. The left column should contain a bar chart showing the average house price by the number of bedrooms (1–2, 3, 4+) for *older* houses, whereas the right column should show the corresponding bars for *newer* houses. Use a shared vertical scale so that prices can be compared across the two panels.

Exercise 4: Inflation and unemployment in the US

In this exercise, you’ll be working with selected macroeconomic variables for the United States reported at monthly frequency obtained from [FRED](#). The data set starts in 1948 and contains observations for a total of 936 months.

1. Load the data from the file `FRED_monthly.csv` located in the `data/FRED` folder. Parse `DATE` as dates and use it as the `DataFrame` index. Sort the observations in chronological order, then print the first 10 observations to get an idea of what the data looks like.
Keep only the columns `Year`, `Month`, `CPI`, and `UNRATE`, retaining the date index. Perform this analysis only on observations prior to 1970 and drop the rest, leaving the sample period 1948–1969.
2. The column `CPI` stores the consumer price index (CPI) for the US. Create a new column `Inflation` which contains *year-over-year inflation in percent*: the percent change in the CPI relative to the same month in the previous year. Apply `pct_change()` to the column `CPI`.

Hints:

- Since this is monthly data, pass `periods=12` to `pct_change()` to get year-over-year changes.
 - Multiply the values returned by `pct_change()` by 100 to get percent values.
3. Compute the average unemployment rate (column `UNRATE`) over the restricted 1948–1969 sample. Create a new column `UNRATE_HIGH` that contains an indicator whenever the unemployment rate is above its average value (“*high unemployment period*”).
 - How many observations fall into the high- and low-unemployment periods?
 - What is the average unemployment rate in the high- and low-unemployment periods?

4. Compute the average year-over-year inflation rate for high- and low-unemployment periods using the available observations. Is there any difference?

Visualize this relationship by creating a scatter plot with the unemployment rate on the x -axis and inflation on the y -axis, and assign different colors to observations from the low- and high-unemployment periods.

5. Use `resample()` to compute the calendar-year average of the monthly year-over-year inflation rates.
 - Report the three years with the highest averages.
 - Create a plot that shows these calendar-year averages over the sample period.

Hint: Use the resampling rule 'YE' when calling `resample()`.